

Tensorboard工具与模型优化

七月在线 张雨石

2017年8月13日

<http://blog.csdn.net/stdcoutzyx>

目录

- Tensorboard 简介
- Tensorboard 使用
- 调参手段
- 调优案例分析

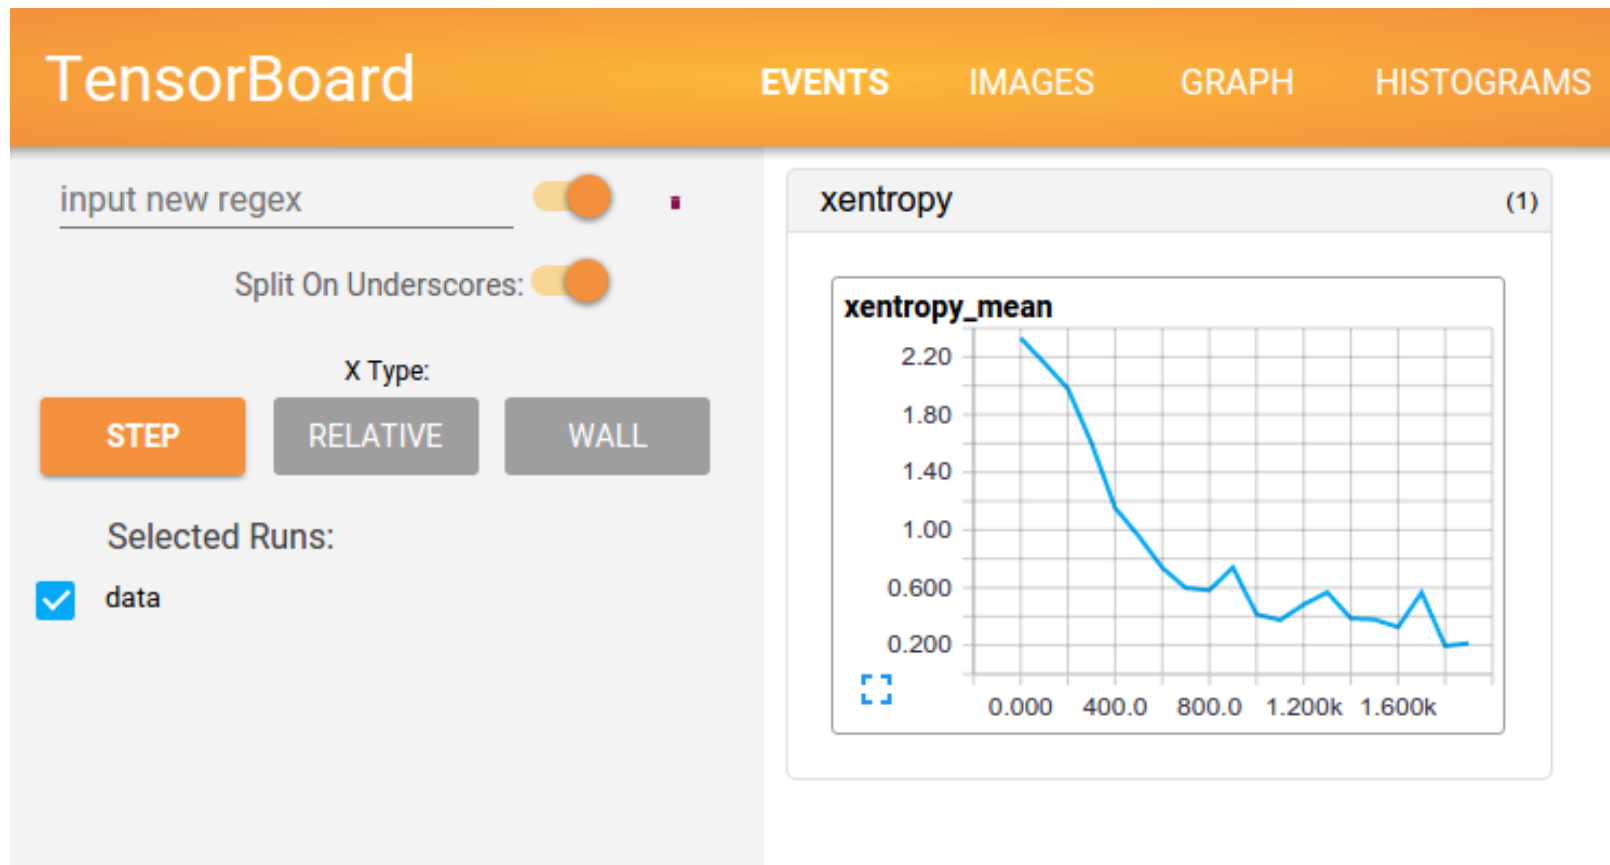


Tensorboard简介

Deep Learning model == Black box



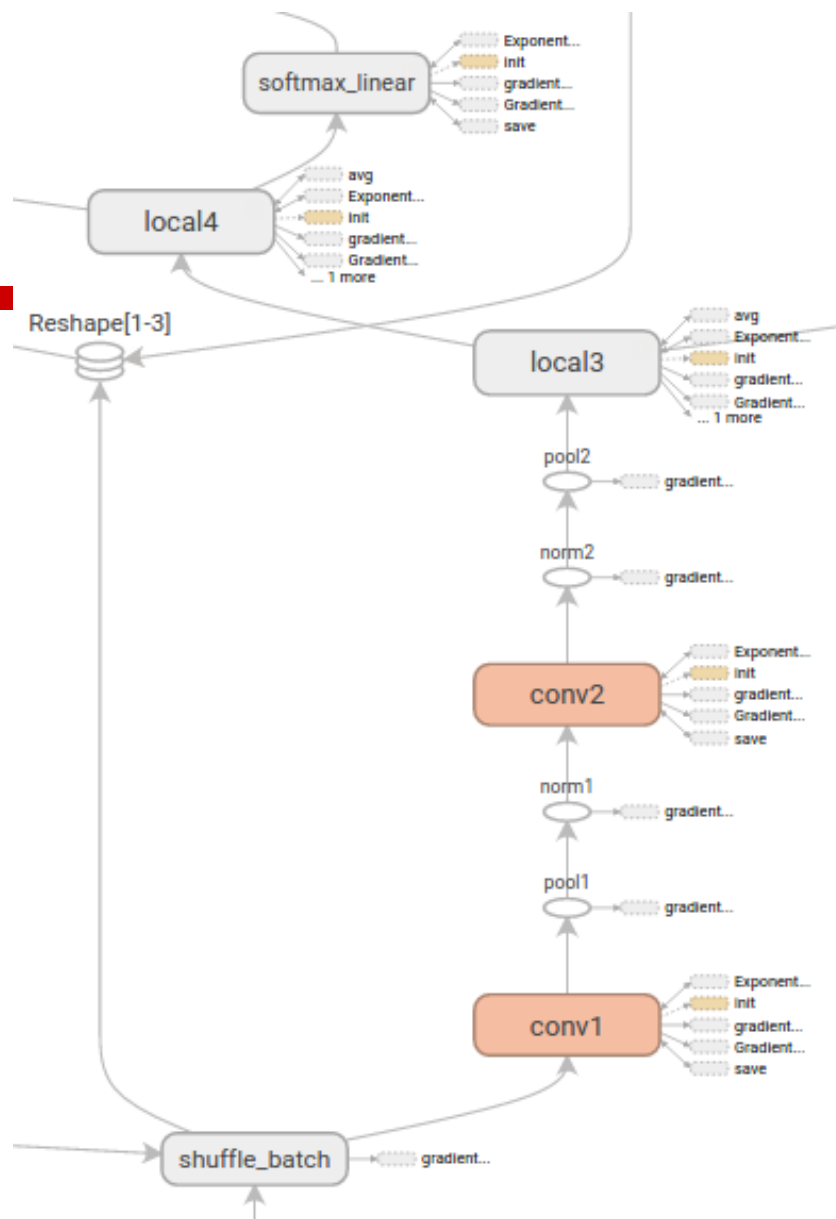
Tensorboard简介



Tensorboard简介

□ 可视化

■ 模型结构




Tensorboard简介

□ 可视化

■ 模型结构


■ 参数分布



Tensorboard简介

□ 可视化

- 模型结构
- 参数分布
- 训练过程



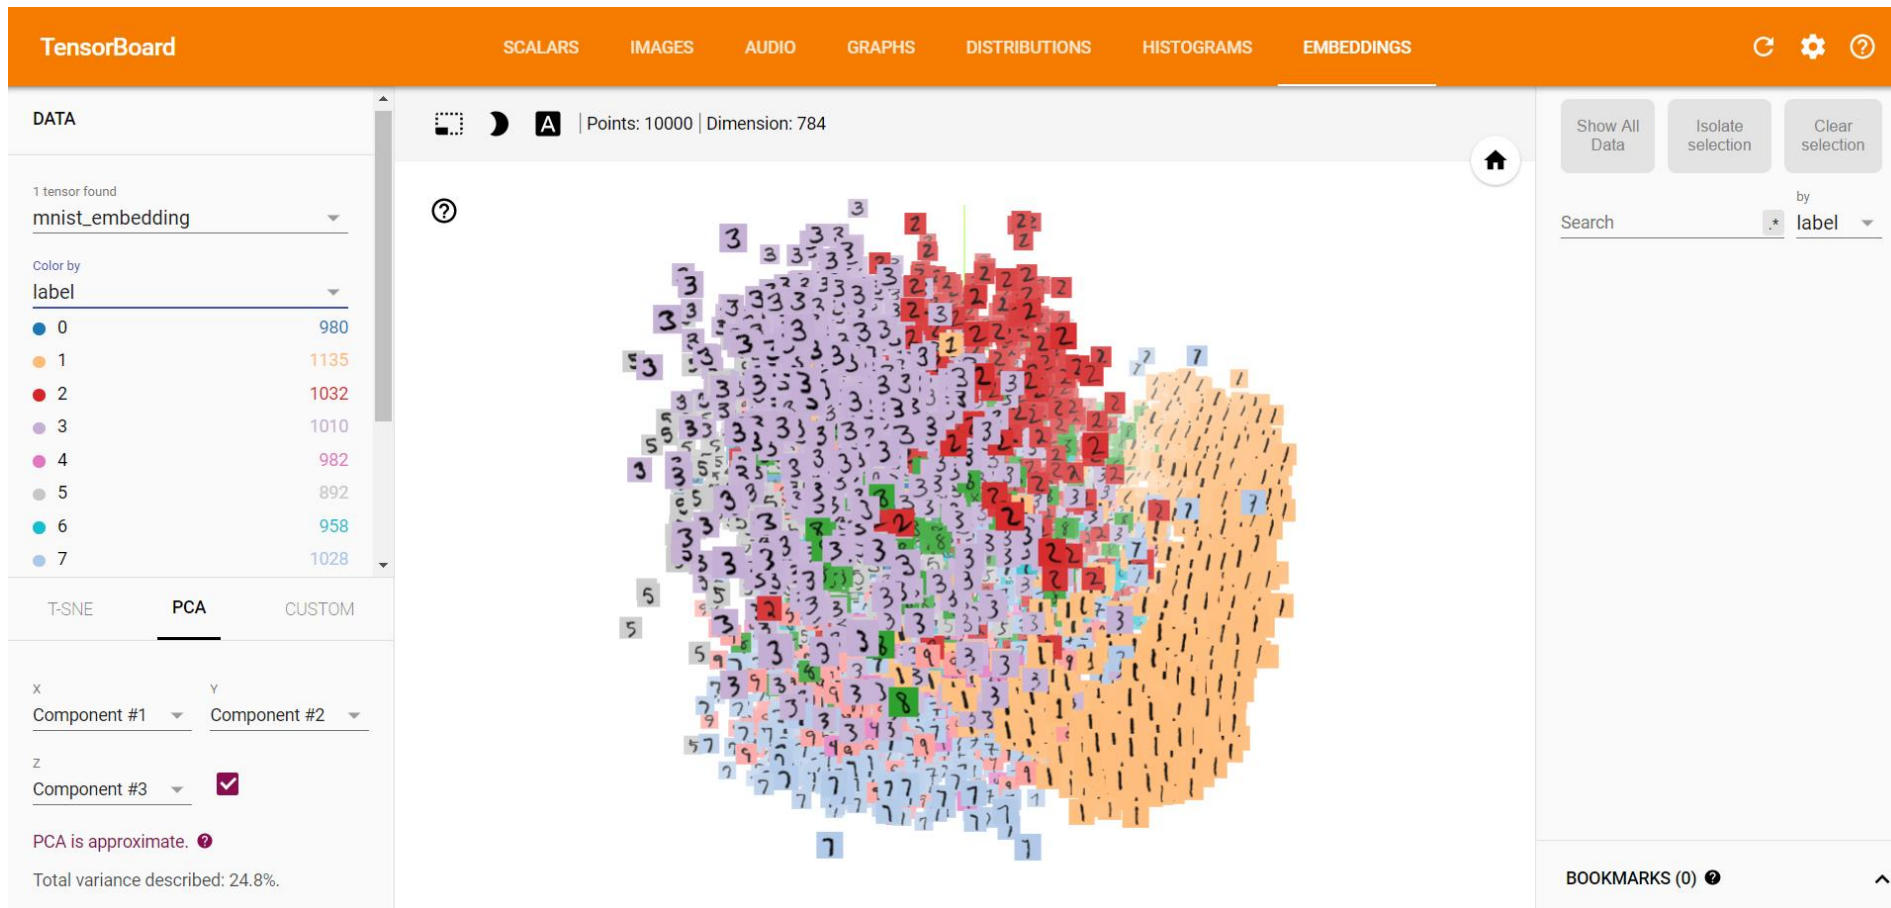
Tensorboard简介

□ 可视化

- 模型结构
- 参数分布
- 训练过程
- Embedding数据分析

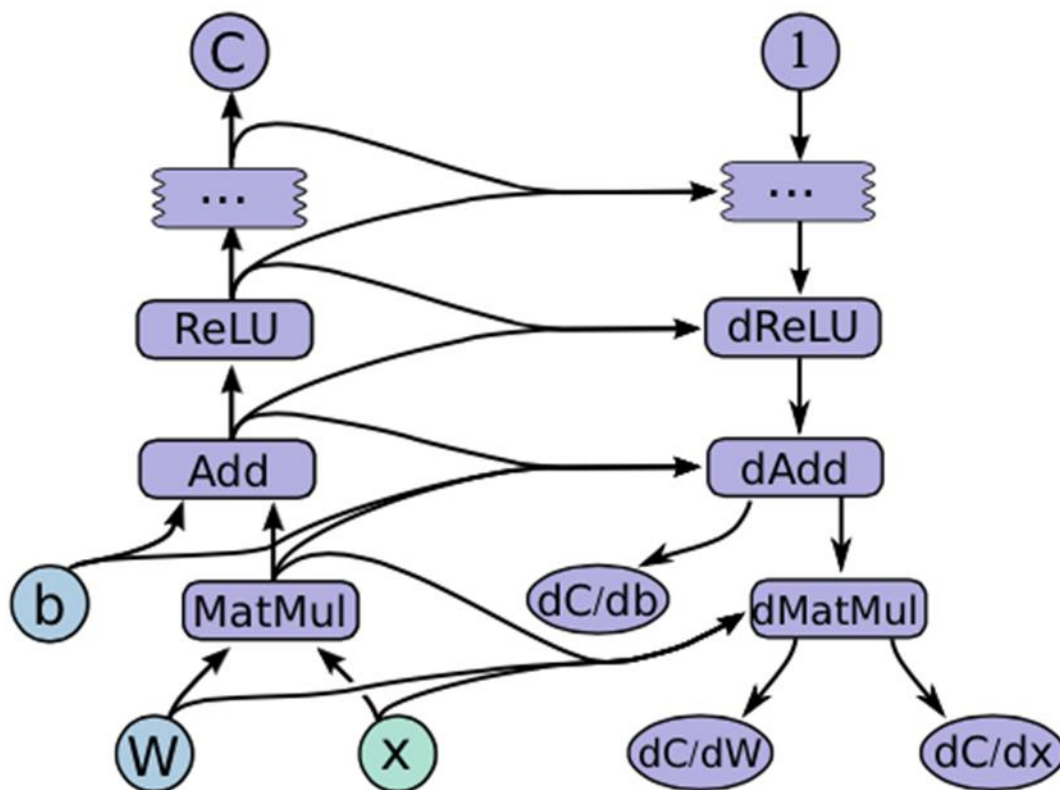


Tensorboard简介



Tensorboard使用-Big Picture

□ Tensorflow 概念回顾



Tensorboard使用-Big Picture

□ Summary_ops

- `tf.summary.tensor_summary`
- `tf.summary.scalar`
- `tf.summary.histogram`
- `tf.summary.audio`
- `tf.summary.image`
- `tf.summary.merge`
- `tf.summary.merge_all`



Tensorboard使用-Big Picture

□ Summary_ops

- 数据序列化
- 写到硬盘/发给tensorboard

□ 其他API

- `tf.summary.FileWriter`
- `tf.summary.FileWriterCache`



Tensorboard使用-Big Picture

□ 回顾：可以做什么

- 输出模型结构
- 跟踪训练状态
- 分析参数
- 分析数据



Tensorboard使用-base code

```
11 def conv2d(x, output_dim, k_h=5, k_w=5, s_h=1, s_w=1, stddev=0.02, name="conv2d"):
12     with tf.variable_scope(name):
13         w = tf.get_variable('w', [k_h, k_w, x.get_shape()[-1], output_dim],
14                             initializer = tf.truncated_normal_initializer(stddev=stddev))
15         conv = tf.nn.conv2d(x, w, strides=[1, s_h, s_w, 1], padding="SAME")
16         biases = tf.get_variable('biases', [output_dim],
17                                 initializer = tf.constant_initializer(0.0))
18         conv = conv + biases
19         return conv
20
21 def max_pool_2d(x, k_h=2, k_w=2, s_h=2, s_w=2, name="max_pool"):
22     return tf.nn.max_pool(
23         x,
24         ksize=[1, k_h, k_w, 1],
25         strides=[1, s_h, s_w, 1],
26         padding="SAME",
27         name=name)
```



Tensorboard使用-base code

```
--
29 def dense(x, output_dim, stddev=0.02, name="dense"):
30     with tf.variable_scope(name):
31         w = tf.get_variable('w', [x.get_shape()[-1], output_dim],
32                             initializer = tf.truncated_normal_initializer(stddev=stddev))
33         biases = tf.get_variable('biases', [output_dim],
34                                 initializer = tf.constant_initializer(0.0))
35         return tf.matmul(x, w) + biases
36
37 def lrelu(x, name="lrelu"):
38     return tf.maximum(x, 0.2 * x)
```



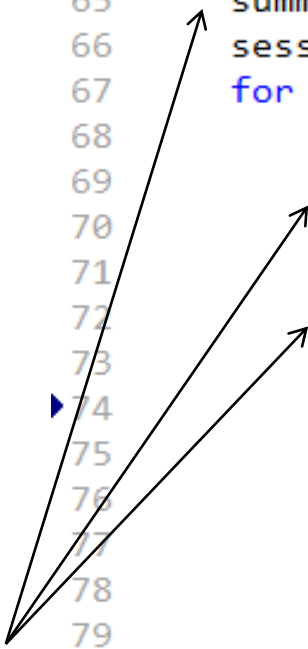
Tensorboard使用-base code

```
40 x = tf.placeholder(tf.float32, [None, 784])
41 x_image = tf.reshape(x, [-1, 28, 28, 1])
42 conv1 = lrelu(conv2d(x_image, 32, name="conv1"))
43 pool1 = max_pool_2d(conv1, name="pool1")
44 conv2 = lrelu(conv2d(pool1, 64, name="conv2"))
45 pool2 = max_pool_2d(conv2, name="pool2")
46 pool2_flatten = tf.reshape(pool2, [-1, 7*7*64])
47 fc1 = lrelu(dense(pool2_flatten, 1024, name="fc1"))
48 dropout_fc1 = tf.nn.dropout(fc1, 0.5)
49 logits = dense(dropout_fc1, 10, name="fc2")
```



Tensorboard使用-模型结构

```
62 init = tf.global_variables_initializer()
63
64 with tf.Session() as sess:
65     summary_writer = tf.summary.FileWriter(FLAGS.ckp_dir, sess.graph)
66     sess.run(init)
67     for i in range(100):
68         batch_xs, batch_ys = mnist.train.next_batch(50)
69         cross_entropy_val, _, summary_str = sess.run(
70             [cross_entropy, train_step, merged_summary],
71             feed_dict={x: batch_xs, y_: batch_ys})
72         summary_writer.add_summary(summary_str, i)
73         if i % 100 == 0:
74             accuracy_val = sess.run(
75                 accuracy,
76                 feed_dict={x: mnist.test.images[0:1000], y_: mnist.test.
77             print "Epoch %4d: cross_entropy: %4.8f, accuracy: %4.8f" % (i, c
78         else:
79             print "Epoch %4d: cross_entropy: %4.8f" % (i, cross_entropy_val)
80
```



Tensorboard使用-模型结构

☐ Launch Tensorboard

- Tensorboard -logdir=dir_path

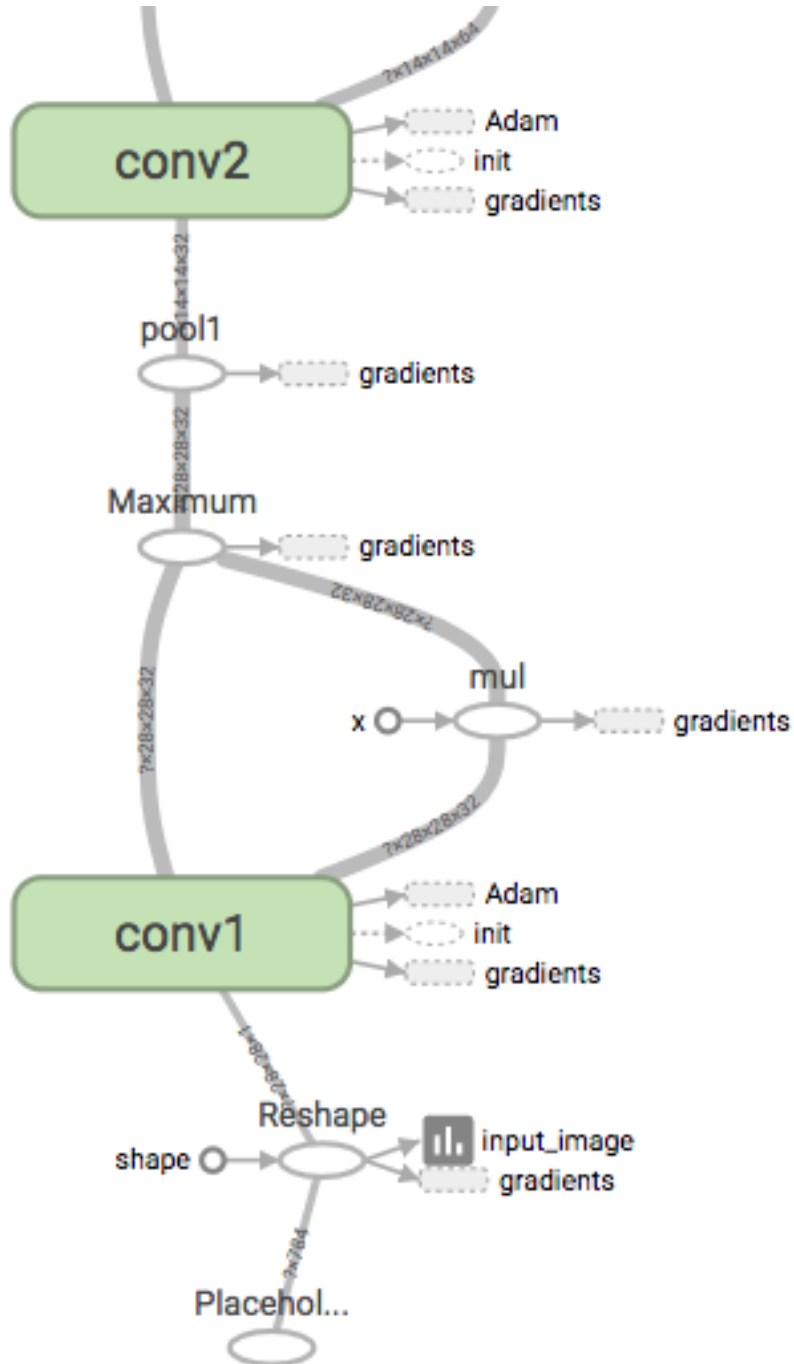
☐ Dir_path

- 单次训练dir
- 多次训练同级dir的父目录
- 两个不相干的dir

- ☐ --logdir=name1:path1,name2:path2



Ten



Tensorboard使用-模型结构

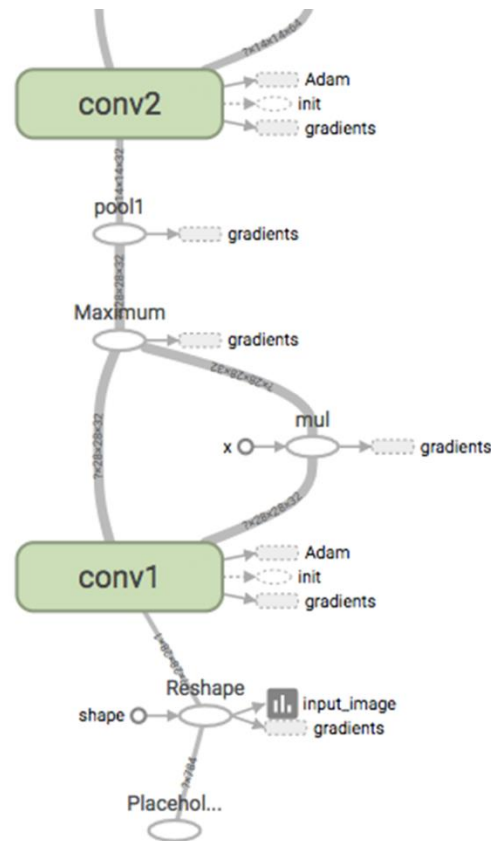
□ 图很乱



□ 命名空间

■ `tf.name_scope`

■ `tf.variable_scope`



Tensorboard-模型结构

□ 命名空间

■ tf.variable_scope

```
with tf.variable_scope("foo"):
    with tf.variable_scope("bar"):
        v = tf.get_variable("v", [1])
        assert v.name == "foo/bar/v:0"
```

```
with tf.variable_scope("foo"):
    v = tf.get_variable("v", [1])
with tf.variable_scope("foo", reuse=True):
    v1 = tf.get_variable("v", [1])
assert v1 == v
```

```
with tf.variable_scope("foo") as scope:
    v = tf.get_variable("v", [1])
    scope.reuse_variables()
    v1 = tf.get_variable("v", [1])
assert v1 == v
```



Tensorboard使用-模型结构


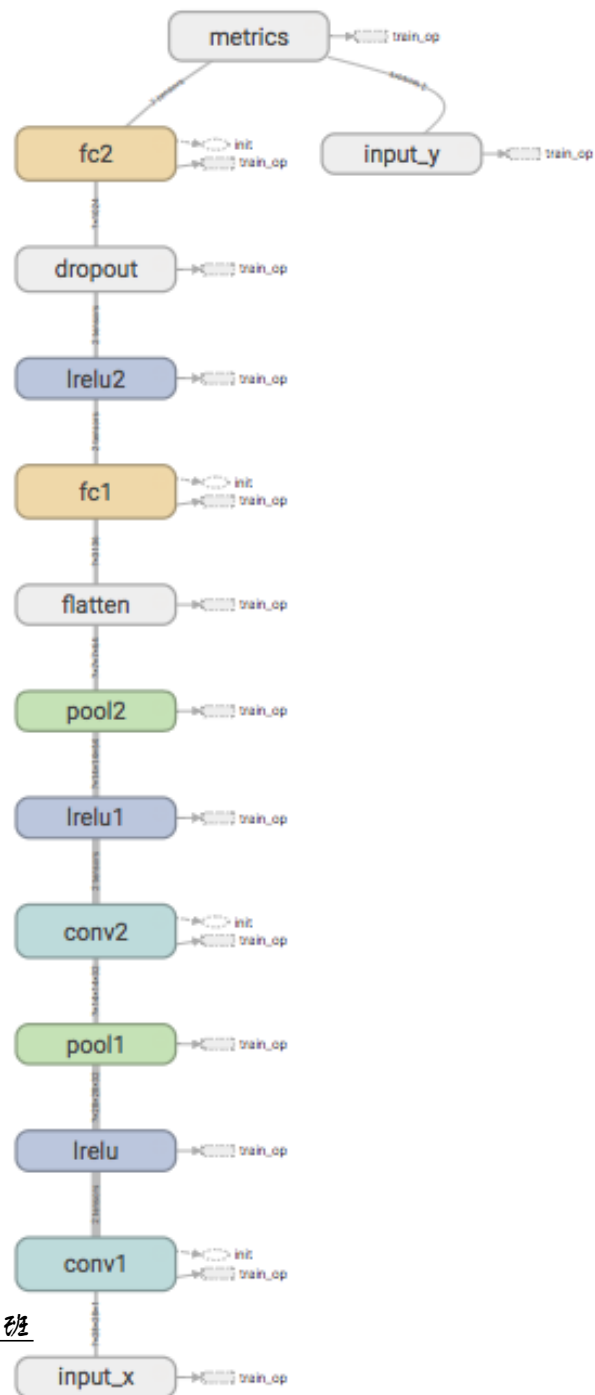
```
21 def max_pool_2d(x, k_h=2, k_w=2, s_h=2, s_w=2, name="max_pool"):  
22     with tf.name_scope(name):  
23         return tf.nn.max_pool(  
24             x,  
25             ksize=[1, k_h, k_w, 1],  
26             strides=[1, s_h, s_w, 1],  
27             padding="SAME",  
28             name='pool')  
29  
30 def dense(x, output_dim, stddev=0.02, name="dense"):  
31     with tf.variable_scope(name):  
32         w = tf.get_variable('w', [x.get_shape()[-1], output_dim],  
33                             initializer = tf.truncated_normal_initializer(stddev=stddev))  
34         biases = tf.get_variable('biases', [output_dim],  
35                                 initializer = tf.constant_initializer(0.0))  
36         return tf.matmul(x, w) + biases  
37  
38 def lrelu(x, name="lrelu"):  
39     with tf.name_scope(name):  
40         return tf.maximum(x, 0.2 * x)
```

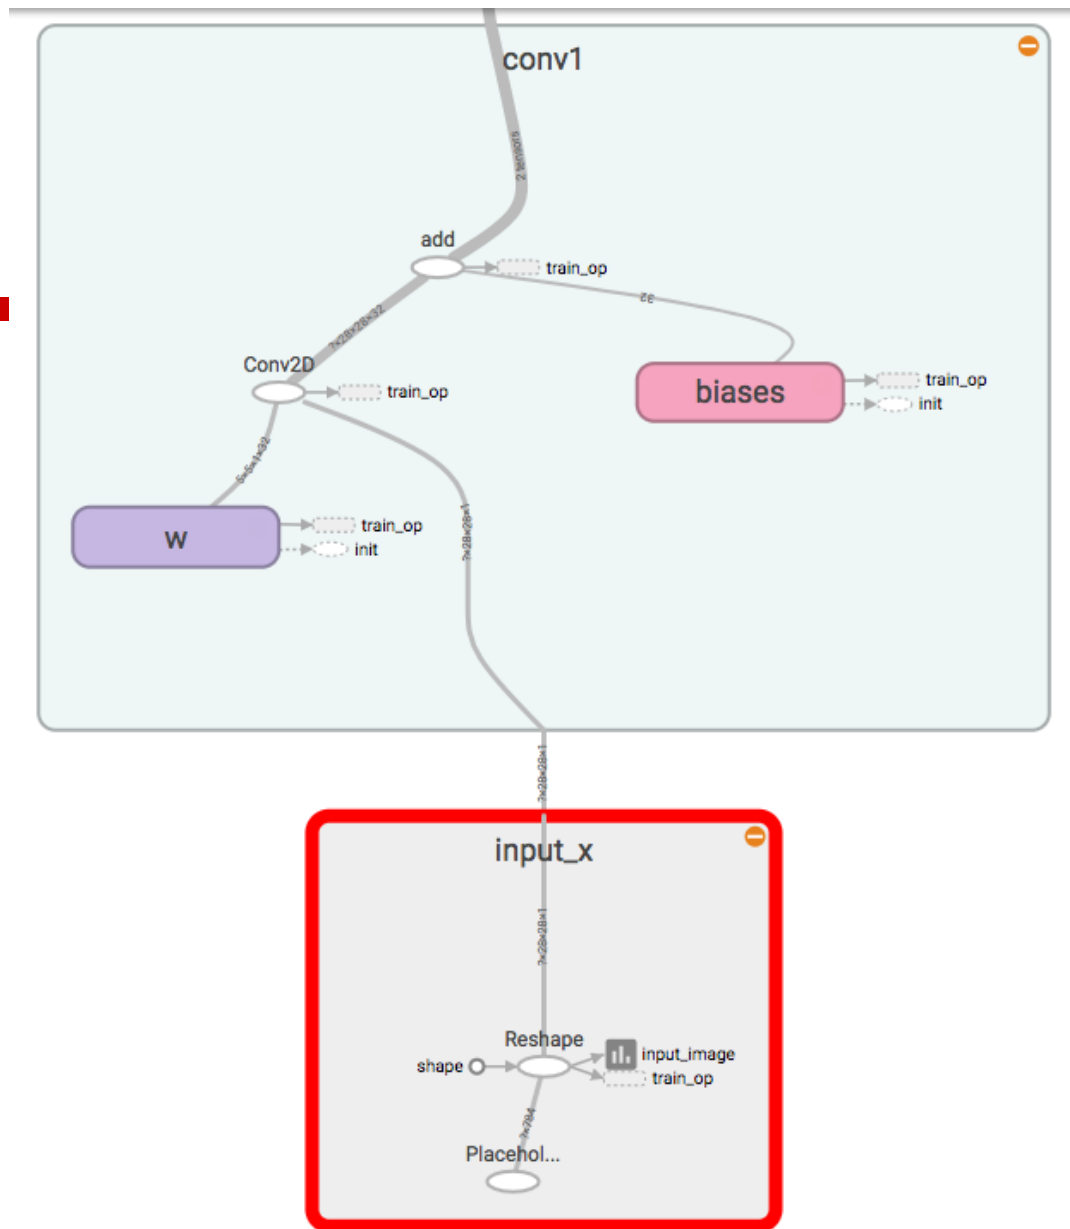


Tensorboard使用-模型结构


```
42 with tf.name_scope("input_x"):
43     x = tf.placeholder(tf.float32, [None, 784])
44     x_image = tf.reshape(x, [-1, 28, 28, 1])
45 conv1 = lrelu(conv2d(x_image, 32, name="conv1"))
46 pool1 = max_pool_2d(conv1, name="pool1")
47 conv2 = lrelu(conv2d(pool1, 64, name="conv2"), name='lrelu1')
48 pool2 = max_pool_2d(conv2, name="pool2")
49 with tf.name_scope("flatten"):
50     pool2_flatten = tf.reshape(pool2, [-1, 7*7*64])
51 fc1 = lrelu(dense(pool2_flatten, 1024, name="fc1"), name='lrelu2')
52 dropout_fc1 = tf.nn.dropout(fc1, 0.5)
53 logits = dense(dropout_fc1, 10, name="fc2")
54
55 with tf.name_scope("input_y"):
56     y_ = tf.placeholder(tf.float32, [None, 10])
57 with tf.name_scope("metrics"):
58     cross_entropy = tf.reduce_mean(
59         tf.nn.sigmoid_cross_entropy_with_logits(logits=logits, labels=y_))
▶ 60     correct_prediction = tf.equal(tf.argmax(logits,1), tf.argmax(y_,1))
61     accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
```







Tensorboard使用-模型结构



train_op
Subgraph: 253 nodes

Attributes (0)

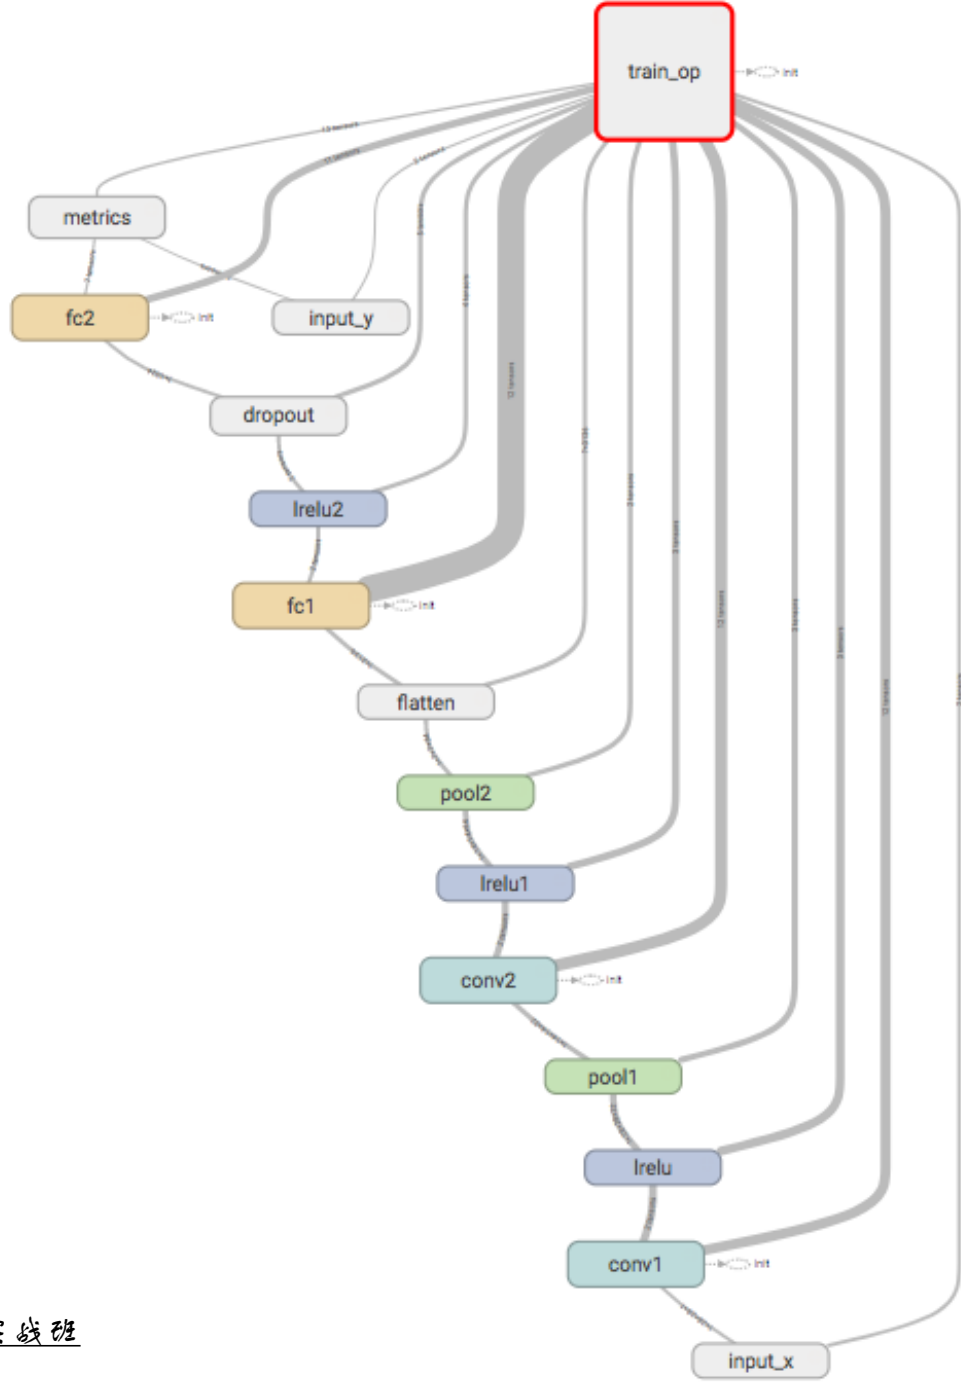
Inputs (14)

	metrics	13 tensors
	fc2	11 tensors
	input_y	2 tensors
	dropout	5 tensors
	lrelu2	4 tensors
	fc1	12 tensors
	flatten/Reshape	?x3136
	pool2	2 tensors
	lrelu1	3 tensors
	conv2	12 tensors
	pool1	3 tensors
	lrelu	3 tensors
	conv1	12 tensors
	input_x	2 tensors

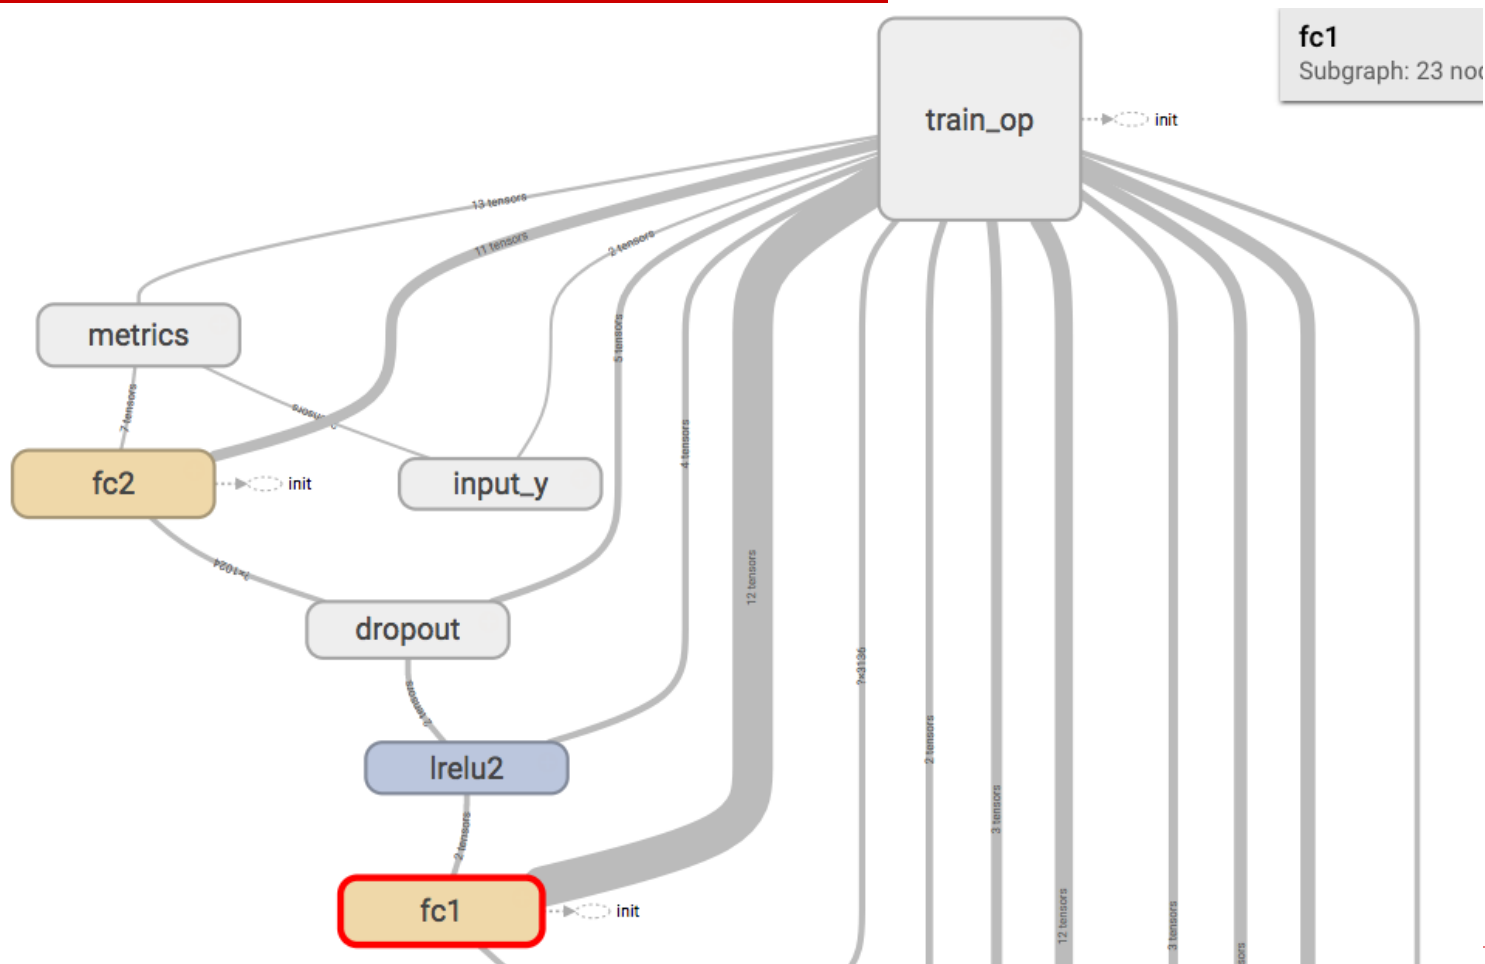
Outputs (1)
✓ Control dependencies

Add to main graph














Tensorboard使用-模型结构



Tensorboard使用-模型结构

Symbol	Meaning
	High-level node representing a name scope. Double-click to expand a high-level node.
	Sequence of numbered nodes that are not connected to each other.
	Sequence of numbered nodes that are connected to each other.
	An individual operation node.
	A constant.
	A summary node.
	Edge showing the data flow between operations.
	Edge showing the control dependency between operations.
	A reference edge showing that the outgoing operation node can mutate the incoming tensor.



Tensorboard使用-模型结构

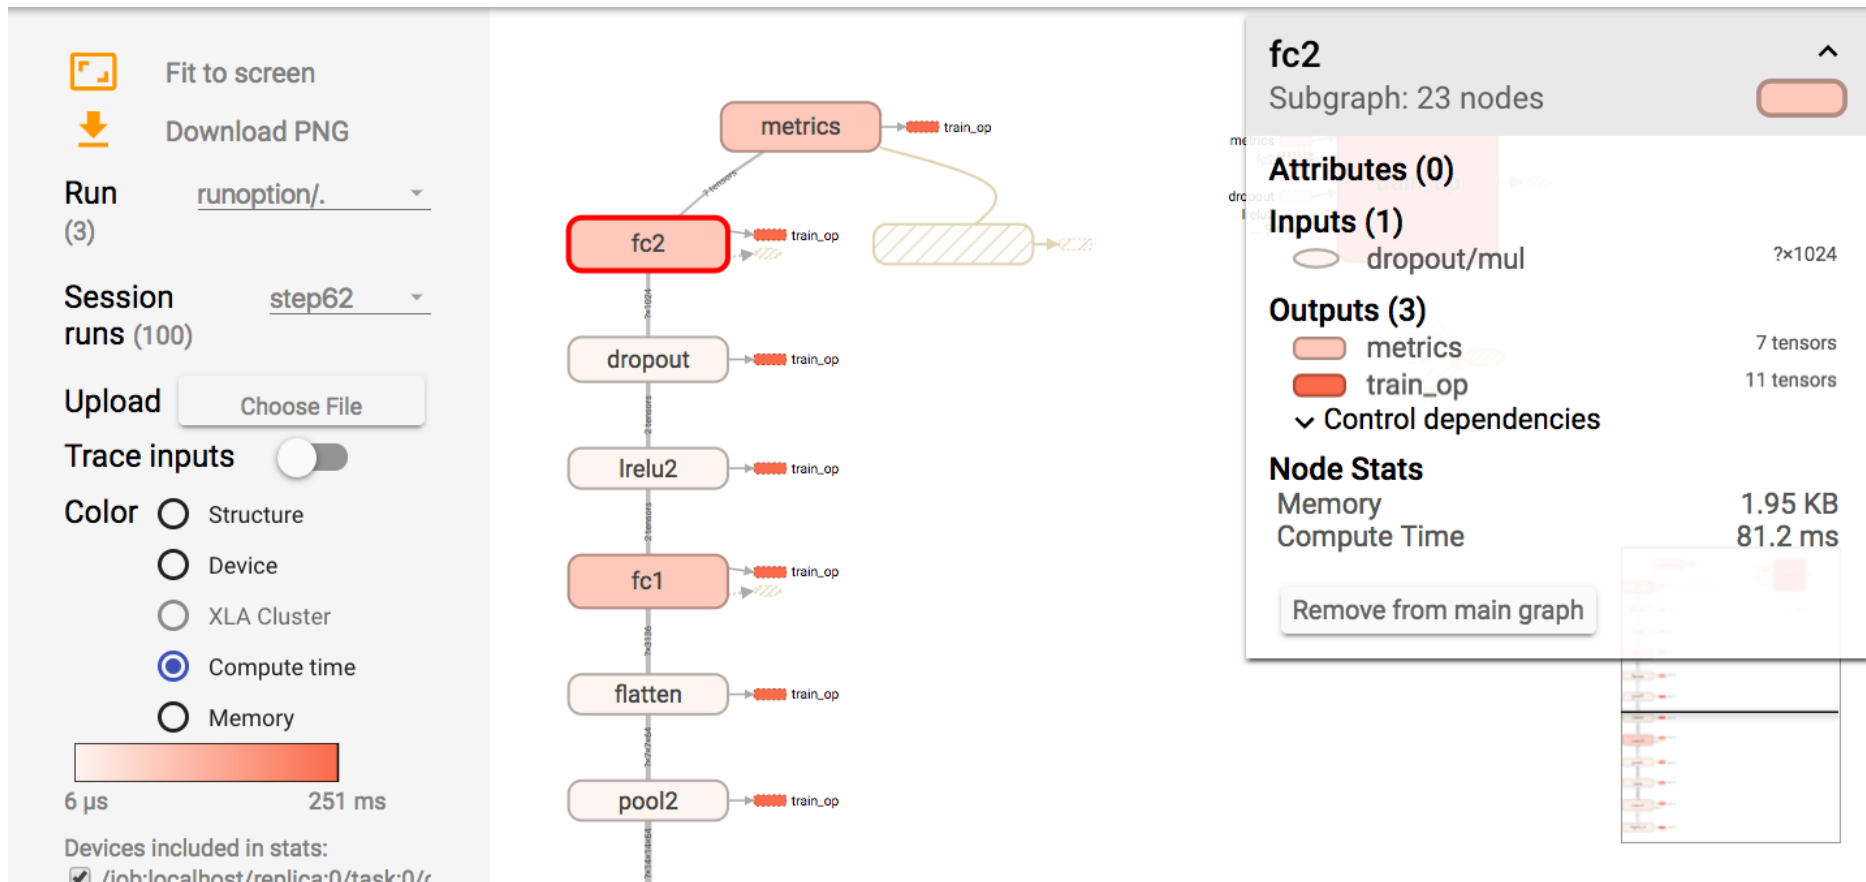
- 就这样？
- System trace
 - Memory
 - Compute time




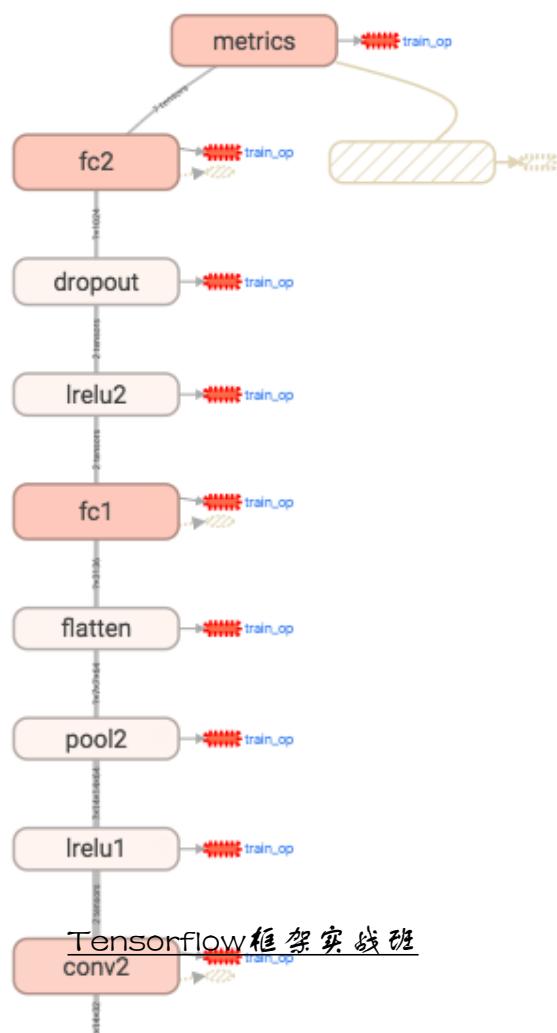
Tensorboard使用-模型结构

```
71 with tf.Session() as sess:
72     summary_writer = tf.summary.FileWriter(FLAGS.ckp_dir, sess.graph)
73     sess.run(init)
74     for i in range(100):
75         batch_xs, batch_ys = mnist.train.next_batch(50)
76         run_options = tf.RunOptions(trace_level=tf.RunOptions.FULL_TRACE)
77         run_metadata = tf.RunMetadata()
78         cross_entropy_val, _, summary_str = sess.run(
79             [cross_entropy, train_step, merged_summary],
80             feed_dict={x: batch_xs, y_: batch_ys},
81             options=run_options,
82             run_metadata=run_metadata)
83         summary_writer.add_run_metadata(run_metadata, 'step%d' % i)
84         summary_writer.add_summary(summary_str, i)
85         if i % 100 == 0:
86             accuracy_val = sess.run(
87                 accuracy,
88                 feed_dict={x: mnist.test.images[0:1000], y_: mnist.test.l
89             print "Epoch %4d: cross_entropy: %4.8f, accuracy: %4.8f" % (i, cr
90         else:
91             print "Epoch %4d: cross_entropy: %4.8f" % (i, cross_entropy_val)
```


Tensorboard使用-模型结构



Tensorboard使用-模型结构



train_op

Subgraph: 253 nodes

Attributes (0)

Inputs (14)

metrics	13 tensors
fc2	11 tensors
input_y	2 tensors
dropout	5 tensors
relu2	4 tensors
fc1	12 tensors
flatten/Reshape	?x3136
pool2	2 tensors
relu1	3 tensors
conv2	12 tensors
pool1	3 tensors
relu	3 tensors
conv1	12 tensors
input_x	2 tensors

Outputs (1)

✓ Control dependencies

Node Stats

Memory

121 MB

Compute Time

246 ms

Add to main graph

julyedu.com

Tensorboard使用-模型结构

conv1
Subgraph: 23 nodes

Attributes (0)

Inputs (1)
input_x/Reshape ?x28x28x1

Outputs (3)
lrelu 2 tensors
train_op 12 tensors
Control dependencies

Node Stats
Memory 9.57 MB
Compute Time 12.4 ms

Remove from main graph

conv2
Subgraph: 23 nodes

Attributes (0)

Inputs (1)
pool1/pool ?x14x14x32

Outputs (3)
lrelu1 2 tensors
train_op 12 tensors
Control dependencies

Node Stats
Memory 2.39 MB
Compute Time 60.4 ms

Remove from main graph



Tensorboard使用-模型结构

fc1 ^
Subgraph: 23 nodes

Attributes (0)

Inputs (1)
○ flatten/Reshape ?×3136

Outputs (3)
○ lrelu 2 tensors
■ train_op 12 tensors
✓ Control dependencies

Node Stats
Memory 200 KB
Compute Time 77.3 ms

Remove from main graph

fc2 ^
Subgraph: 23 nodes

Attributes (0)

Inputs (1)
○ dropout/mul ?×1024

Outputs (3)
■ metrics 7 tensors
■ train_op 11 tensors
✓ Control dependencies

Node Stats
Memory 1.95 KB
Compute Time 78.1 ms

Remove from main graph



Tensorboard使用-模型结构

lrelu1

Subgraph: 2 nodes

Attributes (0)

Inputs (1)

- conv1 2 tensors

Outputs (2)

- pool1/pool $? \times 28 \times 28 \times 32$
- train_op 3 tensors

Node Stats

Memory	9.57 MB
Compute Time	4.61 ms

Remove from main graph

lrelu2

Subgraph: 2 nodes

Attributes (0)

Inputs (1)

- conv2 2 tensors

Outputs (2)

- pool2/pool $? \times 14 \times 14 \times 64$
- train_op 3 tensors

Node Stats

Memory	4.79 MB
Compute Time	3.12 ms

Remove from main graph

lrelu3

Subgraph: 2 nodes

Attributes (0)

Inputs (1)

- fc1 2 tensors

Outputs (2)

- dropout 2 tensors
- train_op 4 tensors

Node Stats

Memory	400 KB
Compute Time	101 μ s

Remove from main graph



Tensorboard使用-训练状态

□ 回顾

- `tf.summary.tensor_summary`
- `tf.summary.scalar`
- `tf.summary.histogram`
- `tf.summary.audio`
- `tf.summary.image`
- `tf.summary.merge`
- `tf.summary.merge_all`



Tensorboard使用-训练状态

```
93 with tf.name_scope("metrics"):
94     cross_entropy = tf.reduce_mean(
95         tf.nn.sigmoid_cross_entropy_with_logits(logits=logits, labels=y_)
96     )
97     correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(y_, 1))
98     accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
99 cross_entropy_summary = tf.summary.scalar('cross_entropy', cross_entropy)
100 accuracy_summary = tf.summary.scalar('accuracy', accuracy)
101 tf.summary.image("input_image", x_image)
102 merged_summary = tf.summary.merge_all()
103 merged_summary_test = tf.summary.merge([cross_entropy_summary, accuracy_summary])
```



Tensorboard使用-训练状态

```
train_writer = tf.summary.FileWriter(train_log_dir, sess.graph)
test_writer = tf.summary.FileWriter(test_log_dir)
sess.run(init)
for i in range(100):
    batch_xs, batch_ys = mnist.train.next_batch(50)
    run_options = tf.RunOptions(trace_level=tf.RunOptions.FULL_TRACE)
    run_metadata = tf.RunMetadata()
    cross_entropy_val, _, summary_str = sess.run(
        [cross_entropy, train_step, merged_summary],
        feed_dict={x: batch_xs, y_: batch_ys},
        options=run_options,
        run_metadata=run_metadata)
    train_writer.add_run_metadata(run_metadata, 'step%d' % i)
    train_writer.add_summary(summary_str, i)
    if (i+1) % 10 == 0:
        accuracy_val, summary_str = sess.run(
            [accuracy, merged_summary_test],
            feed_dict={x: mnist.test.images[0:1000], y_: mnist.test.labels[0:1000]})
        test_writer.add_summary(summary_str, i)
        print "Epoch %4d: cross_entropy: %4.8f, accuracy: %4.8f" % (i, cross_entropy_val,
    else:
        print "Epoch %4d: cross_entropy: %4.8f" % (i, cross_entropy_val)
```



Tensorboard使用-训练状态

```
12 def variable_summaries(var):
13     with tf.name_scope('summaries'):
14         mean = tf.reduce_mean(var)
15         with tf.name_scope('stddev'):
16             stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))
17         tf.summary.scalar('mean', mean)
18         tf.summary.scalar('stddev', stddev)
19         tf.summary.scalar('min', tf.reduce_min(var))
20         tf.summary.scalar('max', tf.reduce_max(var))
21         tf.summary.histogram('histogram', var)
```



Tensorboard使用-训练状态

```
33 def conv2d(x, output_dim, k_h=5, k_w=5, s_h=1, s_w=1, stddev=0.02, name="conv2d"):
34     with tf.variable_scope(name):
35         w = tf.get_variable('w', [k_h, k_w, x.get_shape()[-1], output_dim],
36                             initializer = tf.truncated_normal_initializer(stddev=stddev))
37         conv = tf.nn.conv2d(x, w, strides=[1, s_h, s_w, 1], padding="SAME")
38         biases = tf.get_variable('biases', [output_dim],
39                                 initializer = tf.constant_initializer(0.0))
40         conv = conv + biases
41         with tf.name_scope('w'):
42             variable_summaries(w)
43         with tf.name_scope('biases'):
44             variable_summaries(biases)
45         with tf.name_scope('conv'):
46             variable_summaries(conv)
47         return conv
```



Tensorboard使用-训练状态


```
60 def dense(x, output_dim, stddev=0.02, name="dense"):
61     with tf.variable_scope(name):
62         w = tf.get_variable('w', [x.get_shape()[-1], output_dim],
63                             initializer = tf.truncated_normal_initializer(stddev=stddev))
64         biases = tf.get_variable('biases', [output_dim],
65                                 initializer = tf.constant_initializer(0.0))
66         result = tf.matmul(x, w) + biases
67         with tf.name_scope('w'):
68             variable_summaries(w)
69         with tf.name_scope('biases'):
70             variable_summaries(biases)
71         with tf.name_scope('result'):
72             variable_summaries(result)
73         return result

75 def lrelu(x, name="lrelu"):
76     with tf.name_scope(name):
77         result = tf.maximum(x, 0.2 * x)
78         variable_summaries(result)
79         return result
```




Tensorboard使用-训练状态

Accuracy随迭代次数变化图



Tensorboard使用-训练状态


Cross_entropy随迭代次数变化图




参数Distribution图

conv1


conv1/biases_1/summaries/histogram
train



conv1/conv/summaries/histogram
train




conv1/w_1/summaries/histogram
train




conv2


conv2/biases_1/summaries/histogram
train



conv2/conv/summaries/histogram
train




conv2/w_1/summaries/histogram
train




参数histogram

conv1


conv1/biases_1/summaries/histogram
train



conv1/conv/summaries/histogram
train




conv1/w_1/summaries/histogram
train




conv2


conv2/biases_1/summaries/histogram
train




conv2/conv/summaries/histogram
train




conv2/w_1/summaries/histogram
train




conv1/biases_1/summaries/histogram
train




conv1/biases_1/summaries/histogram
train




conv1/biases_1/summaries/max




conv1/biases_1/summaries/mean




conv1/biases_1/summaries/min




conv1/biases_1/summaries/stddev_1




conv1/conv/summaries/max




conv1/conv/summaries/mean




conv1/conv/summaries/min




conv1/conv/summaries/stddev_1




conv1/w_1/summaries/max




conv1/w_1/summaries/mean



conv1/w_1/summaries/min



conv1/w_1/summaries/stddev_1




Tensorboard使用-模型状态

```
23 def feature_map_summary(conv, dim, name):
24     conv_channels = tf.split(conv,
25                               num_or_size_splits=dim, axis=3)
26     with tf.name_scope(name):
27         len_channel = len(conv_channels)
28         for i in range(len_channel):
29             tf.summary.image('channel-%d' % (i),
30                             conv_channels[i], max_outputs=1)
```



Tensorboard使用-模型状态



```
84 conv1 = lrelu(conv2d(x_image, 32, name="conv1"), name='lrelu1')
85 feature_map_summary(conv1, 32, "conv1_feature_map")
86 pool1 = max_pool_2d(conv1, name="pool1")
87 conv2 = lrelu(conv2d(pool1, 64, name="conv2"), name='lrelu2')
88 feature_map_summary(conv2, 64, "conv2_feature_map")
89 pool2 = max_pool_2d(conv2, name="pool2")
90 with tf.name_scope("flatten"):
91     pool2_flatten = tf.reshape(pool2, [-1, 7*7*64])
92 fc1 = lrelu(dense(pool2_flatten, 1024, name="fc1"), name='lrelu3')
93 dropout_fc1 = tf.nn.dropout(fc1, 0.5)
94 logits = dense(dropout_fc1, 10, name="fc2")
```



conv1_feature_map/channel-6/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))



conv1_feature_map/channel-9/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))



conv1_feature_map/channel-7/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))



conv1_feature_map/channel-10/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))




conv1_feature_map/channel-8/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))




conv1_feature_map/channel-11/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))




conv2_feature_map/channel-39/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))




conv2_feature_map/channel-42/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))




conv2_feature_map/channel-40/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))




conv2_feature_map/channel-43/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))



conv2_feature_map/channel-41/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))



conv2_feature_map/channel-44/image
train
step 22 (Sat Aug 12 2017 11:26:37 GMT+0800 (CST))



Tensorboard使用-分析数据

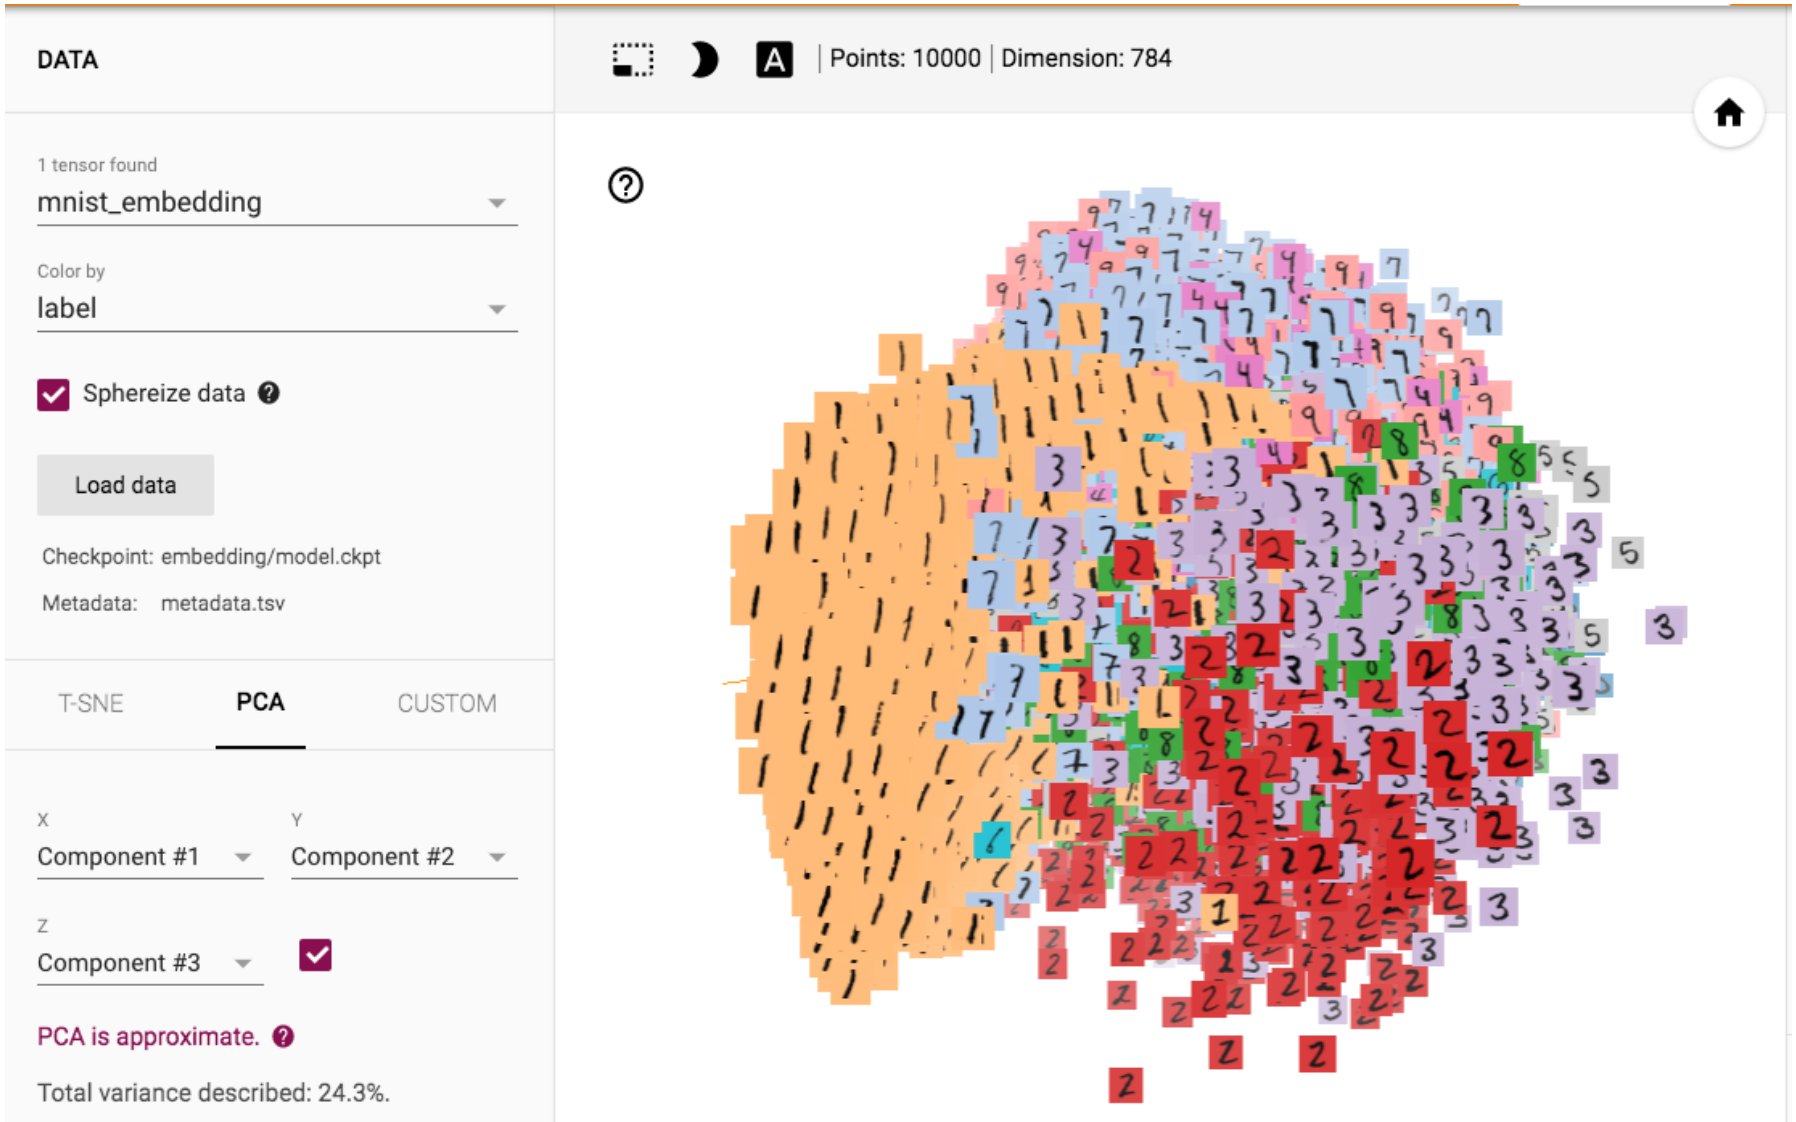
```
21 embedding_var = tf.Variable(mnist.test.images, name="mnist_embedding")
22 saver = tf.train.Saver()
23 sess = tf.Session()
24
25 sess.run(tf.global_variables_initializer())
26 saver.save(sess, os.path.join(path, 'model.ckpt'))
27
28 config = projector.ProjectorConfig()
29
30 embedding = config.embeddings.add()
31 embedding.tensor_name = embedding_var.name
32 embedding.metadata_path = os.path.join(path, 'metadata.tsv')
33 embedding.sprite.image_path = '/Users/zhangyx/workspace/mnist/MNIST_data/mnist_10k_sp
34 embedding.sprite.single_image_dim.extend([28, 28])
35
36 summary_writer = tf.summary.FileWriter(path)
37 projector.visualize_embeddings(summary_writer, config)
```

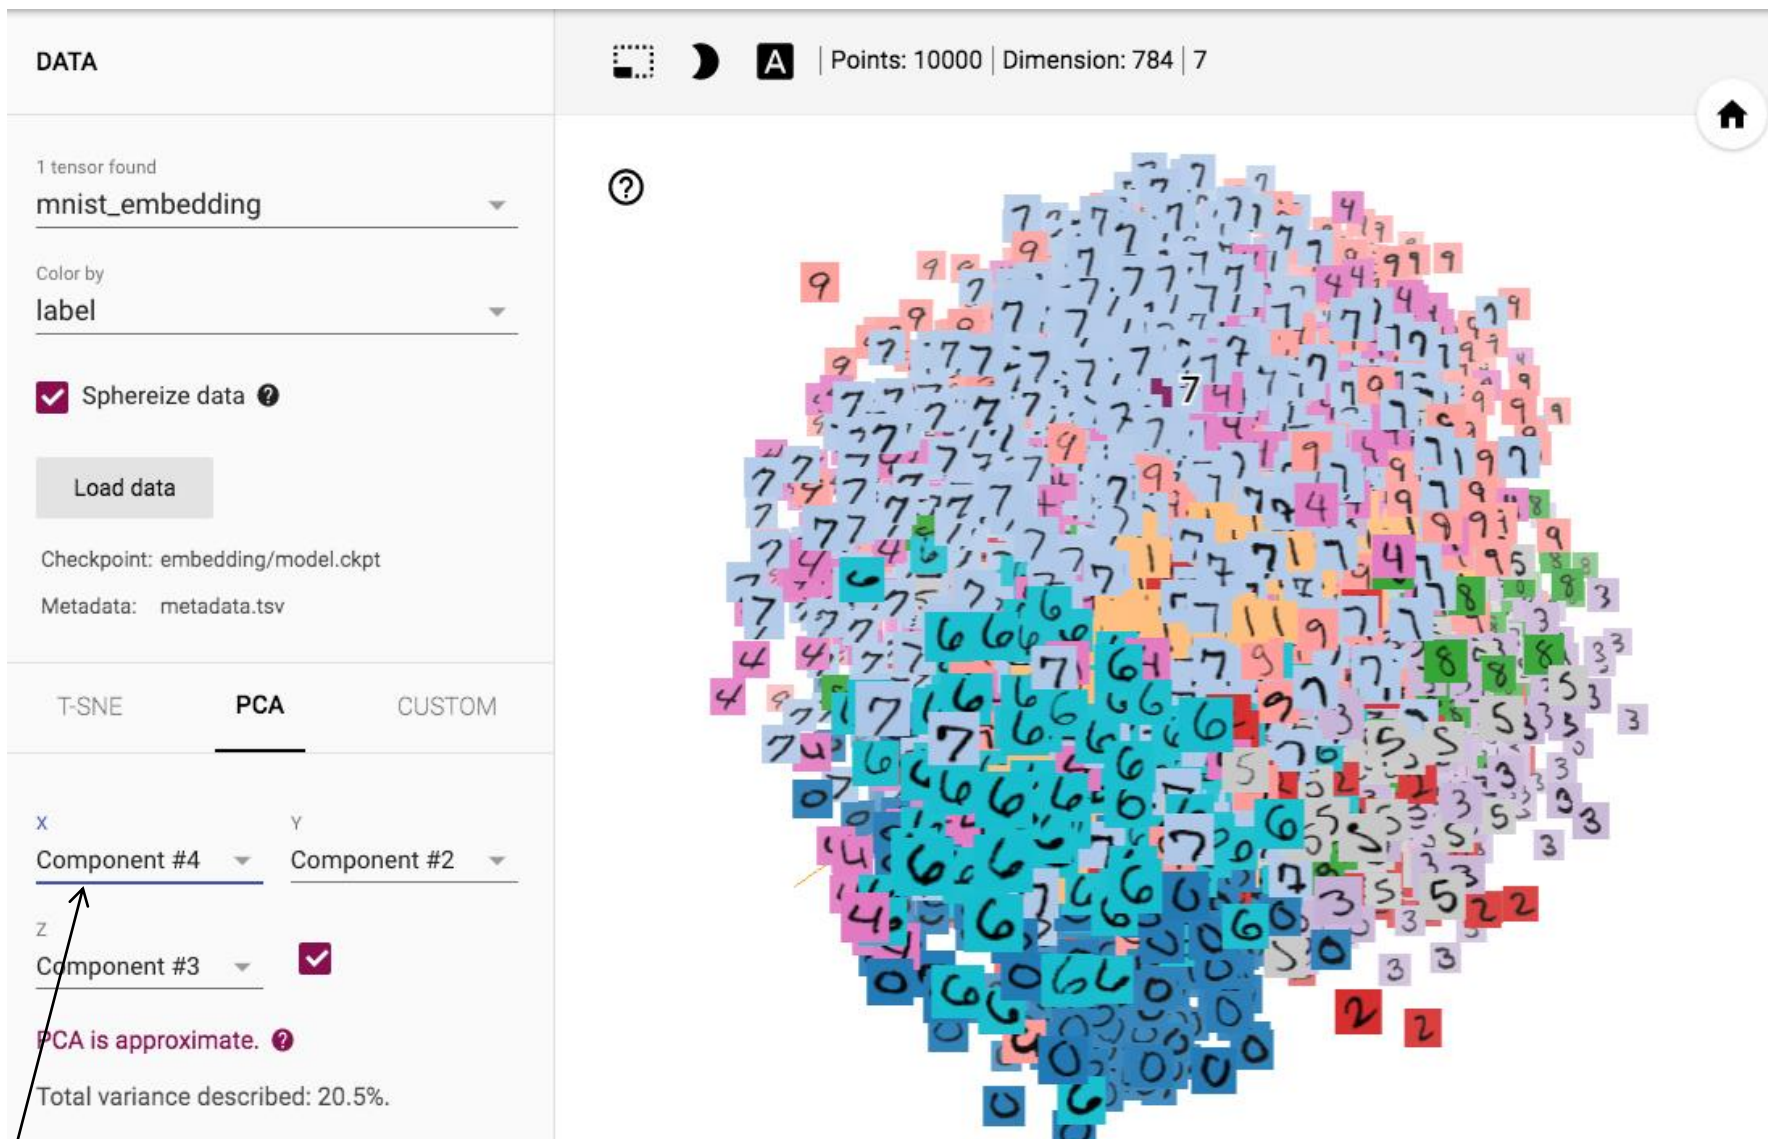


Tensorboard使用-分析数据

6 9 9 6 5 5 3 3 8 1 6 5 6 8 1 9 7 6 8
1 4 0 0 4 1 7 5 7 1 3 3 3 1 6 9 7 4 3
3 4 5 6 7 8 9 7 4 2 0 9 0 1 5 8 8 0 2
5 0 1 2 8 9 1 4 0 9 5 0 8 0 7 7 1 1 2
5 6 7 8 9 3 5 3 2 9 3 2 1 4 5 5 2 3 2
0 9 4 0 1 2 3 4 5 6 7 8 9 0 1 2 3 5 6
0 6 2 5 4 2 3 4 6 0 0 2 0 1 4 5 6 7 8
0 9 6 2 5 4 2 3 4 4 6 0 0 2 0 1 2 3 4
8 4 6 4 9 3 8 4 7 2 5 6 3 6 9 6 3 2 2
2 7 5 7 6 2 9 1 9 0 6 0 6 0 2 0 6 1 5
6 8 4 9 1 7 1 2 3 5 9 6 9 1 1 1 2 9 5
9 0 1 2 3 4 5 6 7 8 0 1 2 3 4 5 6 7 8
2 1 9 3 9 2 0 6 0 4 0 0 1 2 3 4 7 8 9
5 2 9 2 5 8 9 5 0 1 2 4 5 6 0 1 2 3 4
1 7 5 0 1 1 3 8 4 9 1 8 6 8 9 0 1 2 3
9 0 1 2 3 4 7 8 1 3 5 1 7 7 2 1 4 8 3
0 1 2 3 4 5 6 7 8 9 2 1 7 2 5 0 8 0 2
3 9 2 0 6 0 4 0 6 0 1 2 3 4 5 6 7 8 9







DATA

1 tensor found

mnist_embedding

Color by

label

☒ Sphereize data ?

Load data

Checkpoint: embedding/model.ckpt

T-SNE

PCA

CUSTOM

Dimension

2D

3D

Perplexity ?

25

Learning rate ?

10

Re-run

Stop


Iteration: 2986



[How to use t-SNE effectively.](#)



Points: 10000 | Dimension: 784



DATA

1 tensor found

mnist_embedding

Color by

label

☒ Sphereize data ?

Load data

Checkpoint: embedding/model.ckpt

T-SNE

PCA

CUSTOM

Dimension

2D

3D

Perplexity ?

25


Learning rate ?

10

Re-run


Stop

Iteration: 2986

 [How to use t-SNE effectively.](#)



Points: 10000 | Dimension: 784



DATA

1 tensor found

mnist_embedding

Color by

label

☒ Sphereize data ?

Load data

Checkpoint: embedding/model.ckpt

T-SNE

PCA

CUSTOM

Dimension

2D



3D

Perplexity ?



25

Learning rate ?



10

Re-run

Stop

Iteration: 2986



[How to use t-SNE effectively.](#)



Points: 10000 | Dimension: 784



Break!



Tensorboard使用-调优手段

☐ 收集高质量数据

☐ Data Augmentation

- Distortion

- Crop

- Reflection

-

- <https://github.com/mdbloice/Augmentor>

☐ 数据归一化




Tensorboard使用-调优手段

- ☐ `tf.image.flip_left_right`
- ☐ `tf.image.flip_up_down`
- ☐ `tf.image.random_brightness`
- ☐ `tf.image.random_contrast`
- ☐ `tf.image.random_flip_left_right`
- ☐ `tf.image.random_flip_up_down`
- ☐ `tf.image.random_hue`
- ☐ `tf.image.random_saturation`



Tensorboard使用-调优手段



Tensorboard使用-调优手段

☐ 参数初始化

■ 均匀分布初始化 $[-scale, scale]$

☐ Xavier: $Scale = \sqrt{3/n}$

☐ He: $Scale = \sqrt{6/n}$

■ 高斯分布初始化

☐ Xavier: $stddev = \sqrt{1/n}$

☐ He: $stddev = \sqrt{2/n}$

■ N

☐ Xavier: $n = (fan_in + fan_out) / 2$

☐ He: $n = fan_in$



Tensorboard使用-调优手段

□ 参数初始化

- `tf.contrib.keras.initializers.glorot_normal`
- `tf.contrib.keras.initializers.glorot_uniform`
- `tf.contrib.keras.initializers.he_normal`
- `tf.contrib.keras.initializers.he_uniform`
- `tf.contrib.keras.initializers.lecun_uniform`



Tensorboard使用-调优手段

□ 优化方法

- Adam
- Adadelata
- Sgd
- Sgd+momentum



Tensorboard使用-调优手段

☐ 归一化

- 数据归一化

- 梯度归一化

 - ☐ 梯度/mini_size

- Batch Normalization



Tensorboard使用-调优手段

☐ Learning rate

■ 太大

☐ Loss爆炸或者nan

■ 太小

☐ 没有反应

■ 初始设置小的，后期慢慢调大

☐ 1, 0.1, 0.01, 0.001

☐ 梯度截断



Tensorboard使用-调优手段

☐ Batch size

- =1

- 大

- ☐ 吞吐量大

- ☐ 内存计算资源占用多



Tensorboard使用-调优手段

□ Ensemble

- 同样模型结构，不同初始化方法
- 不同的模型结构，选取最好的几个model
- 同样的模型，不同的训练阶段



Tensorboard使用-调优手段

□ 模型结构

- 模型深度
- 每一层的filter数目
- Dropout
- 激活函数



Tensorboard使用-调优手段

□ Early stopping

- Val accuracy一下降就停
- 连续多次val accuracy下降
- 记录最优值，连续多次val accuracy没达到最优值
- 每一段时间对val accuracy计算平均值，平均值下降即停



Tensorboard使用-调优手段

□ 欠拟合

- 增大网络复杂度
- 降低learning rate
- Batch normalization



Tensorboard使用-调优手段

☐ 过拟合

- 丰富数据
- 降低网络复杂度
- 正则化
 - ☐ L1
 - ☐ L2
- Dropout
- Early stopping
- 降低learning rate



Tensorboard使用-Cifar

□ Cifar数据集

■ 32 * 32

■ Trainset

□ 50000

■ Testset

□ 10000

airplane

automobile

bird

cat

deer


dog

frog

horse

ship

truck










who is the best in CIFAR-10 ?



CIFAR-10 49 results collected

Units: accuracy %

Classify 32x32 colour images.

Result	Method	Venue	Details
96.53%	Fractional Max-Pooling 	arXiv 2015	Details
95.59%	Striving for Simplicity: The All Convolutional Net 	ICLR 2015	Details
94.16%	All you need is a good init 	ICLR 2016	Details
94%	Lessons learned from manually classifying CIFAR-10 	unpublished 2011	Details
93.95%	Generalizing Pooling Functions in Convolutional Neural Networks: Mixed, Gated, and Tree 	AISTATS 2016	Details
93.72%	Spatially-sparse convolutional neural networks 	arXiv 2014	
93.63%	Scalable Bayesian Optimization Using Deep Neural Networks 	ICML 2015	



Tensorboard使用-Cifar

```
170 # Randomly crop a [height, width] section of the image.
171 distorted_image = tf.random_crop(reshaped_image, [height, width, 3])
172
173 # Randomly flip the image horizontally.
174 distorted_image = tf.image.random_flip_left_right(distorted_image)
175
176 # Because these operations are not commutative, consider randomizing
177 # the order their operation.
178 # NOTE: since per_image_standardization zeros the mean and makes
179 # the stddev unit, this likely has no effect see tensorflow#1458.
180 distorted_image = tf.image.random_brightness(distorted_image,
181                                              max_delta=63)
182 distorted_image = tf.image.random_contrast(distorted_image,
183                                              lower=0.2, upper=1.8)
184
185 # Subtract off the mean and divide by the variance of the pixels.
186 float_image = tf.image.per_image_standardization(distorted_image)
```



Tensorboard使用-Cifar

```
115 def _variable_with_weight_decay(name, shape, stddev, wd):
116     dtype = tf.float16 if FLAGS.use_fp16 else tf.float32
117     var = _variable_on_cpu(
118         name,
119         shape,
120         tf.truncated_normal_initializer(stddev=stddev, dtype=dtype))
121     if wd is not None:
122         weight_decay = tf.multiply(tf.nn.l2_loss(var), wd, name='weight_loss')
123         tf.add_to_collection('losses', weight_decay)
124     return var

268 labels = tf.cast(labels, tf.int64)
269 cross_entropy = tf.nn.sparse_softmax_cross_entropy_with_logits(
270     labels=labels, logits=logits, name='cross_entropy_per_example')
271 cross_entropy_mean = tf.reduce_mean(cross_entropy, name='cross_entropy')
272 tf.add_to_collection('losses', cross_entropy_mean)
273
274 # The total loss is defined as the cross entropy loss plus all of the weight
275 # decay terms (L2 loss).
276 return tf.add_n(tf.get_collection('losses'), name='total_loss')
```



Tensorboard使用-Cifar

```
188 with tf.variable_scope('conv1') as scope:
189     kernel = _variable_with_weight_decay('weights',
190                                         shape=[5, 5, 3, 64],
191                                         stddev=5e-2,
192                                         wd=0.5)
193     conv = tf.nn.conv2d(images, kernel, [1, 1, 1, 1], padding='SAME')
194     biases = _variable_on_cpu('biases', [64], tf.constant_initializer(0.0))
195     pre_activation = tf.nn.bias_add(conv, biases)
196     conv1 = tf.nn.relu(pre_activation, name=scope.name)
197     _activation_summary(conv1)
198
199 # pool1
200 pool1 = tf.nn.max_pool(conv1, ksize=[1, 3, 3, 1], strides=[1, 2, 2, 1],
201                        padding='SAME', name='pool1')
202
203 # norm1
204 norm1 = tf.nn.lrn(pool1, 4, bias=1.0, alpha=0.001 / 9.0, beta=0.75,
                  name='norm1')
```



Tensorboard使用-Cifar

```
206 # conv2
207 with tf.variable_scope('conv2') as scope:
208     kernel = _variable_with_weight_decay('weights',
209                                         shape=[5, 5, 64, 64],
210                                         stddev=5e-2,
211                                         wd=0.5)
212     conv = tf.nn.conv2d(norm1, kernel, [1, 1, 1, 1], padding='SAME')
213     biases = _variable_on_cpu('biases', [64], tf.constant_initializer(0.1))
214     pre_activation = tf.nn.bias_add(conv, biases)
215     conv2 = tf.nn.relu(pre_activation, name=scope.name)
216     _activation_summary(conv2)
217
218 # norm2
219 norm2 = tf.nn.lrn(conv2, 4, bias=1.0, alpha=0.001 / 9.0, beta=0.75,
220                  name='norm2')
221 # pool2
222 pool2 = tf.nn.max_pool(norm2, ksize=[1, 3, 3, 1],
223                          strides=[1, 2, 2, 1], padding='SAME', name='pool2')
```



```

225 # local3
226 with tf.variable_scope('local3') as scope:
227     # Move everything into depth so we can perform a single matrix multiply.
228     reshape = tf.reshape(pool2, [FLAGS.batch_size, -1])
229     dim = reshape.get_shape()[1].value
230     weights = _variable_with_weight_decay('weights', shape=[dim, 384],
231                                           stddev=0.04, wd=0.004)
232     biases = _variable_on_cpu('biases', [384], tf.constant_initializer(0.1))
233     local3 = tf.nn.relu(tf.matmul(reshape, weights) + biases, name=scope.name)
234     _activation_summary(local3)
235
236 # local4
237 with tf.variable_scope('local4') as scope:
238     weights = _variable_with_weight_decay('weights', shape=[384, 192],
239                                           stddev=0.04, wd=0.004)
240     biases = _variable_on_cpu('biases', [192], tf.constant_initializer(0.1))
241     local4 = tf.nn.relu(tf.matmul(local3, weights) + biases, name=scope.name)
242     _activation_summary(local4)
243
244 with tf.variable_scope('softmax_linear') as scope:
245     weights = _variable_with_weight_decay('weights', [192, NUM_CLASSES],
246                                           stddev=1/192.0, wd=0.0)
247     biases = _variable_on_cpu('biases', [NUM_CLASSES],
248                               tf.constant_initializer(0.0))
249     softmax_linear = tf.add(tf.matmul(local4, weights), biases, name=scope.name)
250     _activation_summary(softmax_linear)

```



Tensorboard使用-Cifar

```
324 num_batches_per_epoch = NUM_EXAMPLES_PER_EPOCH_FOR_TRAIN / FLAGS.batch_size
325 decay_steps = int(num_batches_per_epoch * NUM_EPOCHS_PER_DECAY)
326
327 # Decay the learning rate exponentially based on the number of steps.
328 lr = tf.train.exponential_decay(INITIAL_LEARNING_RATE,
329                                 global_step,
330                                 decay_steps,
331                                 LEARNING_RATE_DECAY_FACTOR,
332                                 staircase=True)
333 tf.summary.scalar('learning_rate', lr)
334
335 # Generate moving averages of all losses and associated summaries.
336 loss_averages_op = _add_loss_summaries(total_loss)
337
338 # Compute gradients.
339 with tf.control_dependencies([loss_averages_op]):
340     opt = tf.train.GradientDescentOptimizer(lr)
341     grads = opt.compute_gradients(total_loss)
```



Tensorboard使用-Cifar

```
343 # Apply gradients.
344 apply_gradient_op = opt.apply_gradients(grads, global_step=global_step)
345
346 # Add histograms for trainable variables.
347 for var in tf.trainable_variables():
348     tf.summary.histogram(var.op.name, var)
349
350 # Add histograms for gradients.
351 for grad, var in grads:
352     if grad is not None:
353         tf.summary.histogram(var.op.name + '/gradients', grad)
354
355 # Track the moving averages of all trainable variables.
356 variable_averages = tf.train.ExponentialMovingAverage(
357     MOVING_AVERAGE_DECAY, global_step)
358 variables_averages_op = variable_averages.apply(tf.trainable_variables())
359
360 with tf.control_dependencies([apply_gradient_op, variables_averages_op]):
361     train_op = tf.no_op(name='train')
```



Tensorboard使用-Cifar

□ 86%

■ 100k steps

■ Batch_size=128

System	Step Time (sec/batch)	Accuracy
<hr/>		
1 Tesla K20m	0.35-0.60	~86% at 60K steps (5 hours)
1 Tesla K40m	0.25-0.35	~86% at 100K steps (4 hours)



Tensorboard使用-Cifar


```
79 def _activation_summary(x):
80     tensor_name = re.sub('%s_[0-9]*/' % TOWER_NAME, '', x.op.name)
81     tf.summary.histogram(tensor_name + '/activations', x)
82     tf.summary.scalar(tensor_name + '/sparsity',
83                       tf.nn.zero_fraction(x))
```



Tensorboard使用-Cifar


accuracy

accuracy




local3

local3/local3/local3/sparsity



local4

local4/local4/local4/sparsity




Tensorboard使用-Cifar快速实验

❑ 错误的转换图片

input_image


input_image/image/0

train
step 4328 (Sun Aug 13 2017 15:17:57 GMT+0800 (CST))




input_image/image/1

train
step 4328 (Sun Aug 13 2017 15:17:57 GMT+0800 (CST))




input_image/image/2

train
step 4328 (Sun Aug 13 2017 15:17:57 GMT+0800 (CST))





Tensorboard使用-Cifar快速实验



Tensorboard使用-Cifar快速实验

□ Batch_size



Tensorboard使用-Cifar


□ Initialize

Write a regex to filter runs

- ✓ he/test
- he/train
- ✓ stddev_0.02/test
- stddev_0.02/train
- ✓ xavier/test
- xavier/train

accuracy

accuracy



Summary

- Tensorboard介绍
- 调优手段
- Cifar10实例



谢谢！

